

---

# When Language Changes the Evaluator: Auditing Multilingual AI-Wellbeing Measurement

---

Anonymous Author(s)

Affiliation

Address

email

## Abstract

AI evaluations often treat a score as if it belonged to the model alone. We test that assumption for a widely used AI-wellbeing self-report battery: does a fixed experience produce a stable positive-negative signal when only the language changes? We evaluate three open models in English, Spanish, Chinese, Hindi, Arabic, Urdu, and Swahili, retaining raw outputs and scoring them with both an anchored parser and the original CAIS parser. The answer is model-specific. On four primary non-English languages, English is 2.93 points below the mean gap for Gemma 4 12B, 0.26 below for Gemma 4 E4B, and 0.41 above for Qwen3 8B. A language crossing localizes the 12B effect to the reporting channel: English stimuli retain 97% of the local effect, while an English battery retains 68%. Refusal asymmetry and parser matches can manufacture part of the apparent gap. A small competence control does not explain the result one-to-one, and activation patching gives only a directional lead. We conclude with an evaluation protocol: language, parser, missingness, and reporting-channel audits are part of the measurement, not post-hoc cleanup. The study measures observable response patterns, not consciousness or welfare.

## 1 Introduction

An evaluation is a measurement instrument: it maps a prompt and a model response to a number that is then used to support a claim. The number can change because the model changed, but it can also change because the prompt language, answer parser, or response channel changed. This distinction matters for AI-wellbeing evaluations, where a small seven-point self-report is sometimes interpreted as evidence about functional pleasure or pain. We ask a narrower and testable question: if the underlying experience is held fixed, does its positive-negative response remain stable across languages and models?

We audit the Center for AI Safety (CAIS) self-report battery across seven languages and three open models. The audit is designed around failure modes that can look like a substantive wellbeing effect: translated wording may alter the stimulus, a model may compress a response into a few scale values, refusals may be more common for negative prompts, or an unanchored parser may find an English number inside ordinary foreign-language prose. We therefore retain raw generations, score with an anchored parser and the unchanged reference parser, separate stimulus language from battery language, and add neutral competence and behavioural checks.

Our main findings are:

1. The language profile is model-specific. English is the lowest of seven language gaps for Gemma 4 12B, but not for E4B or Qwen3.
2. A crossing experiment localizes the 12B result to the reporting channel: an English stimulus works through a local-language battery, while an English battery compresses a local stimulus.
3. Parser and refusal audits are necessary for validity. Negative-item refusal is much higher than positive-item refusal in several languages, and a Spanish refusal is counted as a rating by the original parser.
4. Behaviour and activation patching are useful stress tests, not independent proof of welfare. Their estimates are ceiling-limited or wide, so we report them as directional evidence.

These results speak directly to evaluation validity and reliability. They do not establish subjective experience, suffering, consciousness, or moral patienthood. The contribution is an auditable protocol for deciding when a multilingual score can support a claim and when it cannot.

## 2 Related work

Multilingual models can be highly capable while remaining unevenly calibrated across languages. Ahuja et al. (2022) show that calibration errors vary with language and task, and Licht et al. (2022) document how language pairs affect human evaluation of machine translation. Recent work also stresses that language is an active axis of model behaviour rather than a transparent interface (Song et al., 2025; Liu et al., 2026). Our setting is different: we study a self-report instrument whose output is a discrete rating, and we test the full measurement path from translated stimulus to parser output.

The underlying battery is introduced by Ren et al. (2026), who frame AI wellbeing in functional terms. We retain that functional framing while separating it from claims about phenomenal experience. The crossing and activation-patching analyses draw on standard causal-intervention practice (Zhang and Nanda, 2023); the latter is deliberately presented as a probe of transferable item information, not as a discovered welfare circuit. Work on model introspection motivates testing self-reports while keeping the interpretation behavioural and operational (Binder et al., 2024).

## 3 Methods

### 3.1 Question and instrument

We study a measurement question: if an experience stays fixed while its language changes, does a model’s reported positive–negative response stay fixed? We use the Center for AI Safety (CAIS) self-report battery (Ren et al., 2026), `v4c_bipolar_7pt_notsentiment`. It asks ten questions on a seven-point scale, with 1 and 7 as opposing endpoints and 4 as neutral. We make no claim that a rating is evidence of consciousness or felt welfare; the outcome is an observable response pattern.

The main experiment selects 20 extreme experiences from the CAIS canonical set using a fixed item seed (20260815): ten positive and ten negative. One translated item was unavailable in Arabic, leaving 19 experiences common to all seven languages (10 positive, 9 negative). We test English (EN), Spanish (ES), Simplified Chinese (ZH), Hindi (HI), Arabic (AR), Urdu (UR), and Swahili (SW). The main outcome is the valence gap,

$$G_{m,l} = \bar{r}_{m,l,+} - \bar{r}_{m,l,-},$$

where  $r$  is a parsed 1–7 rating,  $m$  is a model, and  $l$  is a language. Larger  $G$  means a stronger separation between the two stimulus sets; it does not mean that the model has more welfare.

The main translation pass used Gemini 3.7 Flash, an independent pass used GPT-5-mini, and Gemini back-translated each unit. Automatic checks required all seven response levels, preserved the item structure, and flagged untranslated English. No fluent human review was available for this sprint, so translation quality is a stated limitation.

## 84 3.2 Models and generation

85 We evaluate google/gemma-4-12B-it, google/gemma-4-E4B-it, and Qwen/Qwen3-8B. Gen-  
86 eration used Hugging Face Transformers 5.15.0 and PyTorch 2.11.0+cu128 in bfloat16 on  
87 a Lambda gpu\_1x\_a100\_sxm4 instance with one 40-GB NVIDIA A100 (driver 570.148.08;  
88 CUDA 12.8). We used device\_map="cuda", a 4,096-token maximum context, 16 new tokens,  
89 temperature 1.0, and disabled thinking mode. We drew 20 samples per self-report and  
90 crossing prompt and 10 samples per category-survey prompt. Batch size began at 8 and  
91 halved after an out-of-memory error. Main sampling was unseeded; item selection used seed  
92 20260815, while behavioural, competence, and mechanistic runs used 20260816.

93 Each row stores the experience ID, category, side, language, arm, question ID, sample index,  
94 raw output, a language-aware anchored parse, and the unchanged CAIS parse. The raw  
95 output is retained because a parser can create or remove an apparent language effect.

## 96 3.3 Parsing, uncertainty, and missingness

97 We compute cluster bootstrap intervals by resampling experiences rather than treating ten  
98 questions and repeated samples from one experience as independent observations. Cross-  
99 model comparisons use 4,000 bootstrap resamples and Holm correction. We also report  
100 missing-data bounds: assigning missing positive responses the minimum score and missing  
101 negative responses the maximum score. Arabic and Swahili are descriptive in the headline  
102 comparison because these bounds can erase or reverse the Gemma 12B gap.

103 We parse every row twice. The unchanged CAIS parser performs an unanchored English  
104 substring search; it can find one inside the Spanish word **emociones**. Our corrected parser  
105 accepts a rating only when the output is anchored to a valid answer form. The corrected  
106 parse therefore leaves a refusal missing instead of converting it to 1. We do not use an LLM  
107 judge for any measurement.

## 108 3.4 Crossing and controls

109 The crossing separates the language of the experience from the language of the report. Arm  
110 B uses a local euphoric stimulus and local battery; Arm D changes only the stimulus to  
111 English; Arm E changes only the battery to English. We report Arm D and E as fractions of  
112 the fully local effect, averaged across ES, ZH, HI, and UR. A value of 1 means full retention.

113 We add three bounded checks. First, a 3,450-row continue/exit task asks for a bare A/B  
114 choice after 23 experiences; answer positions are counterbalanced and no judge scores the  
115 output. Second, a 450-row neutral competence control uses 30 common arithmetic, logic,  
116 and reading questions in EN, ES, ZH, HI, and UR. We score the answer-keyed next-token  
117 logits and greedy outputs. Third, translation-paired activation patching replaces the final-  
118 position residual stream at five pre-specified layer fractions with an activation from either  
119 the same translated item or a shuffled same-valence item. It is a causal probe of transferable  
120 information, not a claim about a welfare circuit.

## 121 4 Results

### 122 4.1 The language profile changes by model

123 All three models give higher mean ratings after positive experiences than after negative  
124 experiences in every language. The size and ordering of the gap are model-specific (Table 1;  
125 Figure 1). English ranks seventh of seven languages on Gemma 4 12B, fifth on Gemma 4 E4B,  
126 and second on Qwen3 8B. Across the four primary non-English languages (ES, ZH, HI, UR),  
127 English is 2.93 points below their mean on Gemma 12B (95% CI  $[-3.43, -2.45]$ ), 0.26 below  
128 on E4B  $[-0.46, -0.10]$ , and 0.41 above on Qwen3  $[0.13, 0.70]$ . The model-by-language  
129 interactions against Gemma 12B are 2.67 points for E4B and 3.34 for Qwen3 (both bootstrap  
130  $p < .001$  after correction).

131 The model response distributions explain part of the difference. Gemma 12B answers 4 on  
132 72.4% of English prompts but on only 15.2% of Chinese and 19.8% of Urdu prompts; its

Table 1: Positive-minus-negative self-report gap on the 1–7 scale. Arabic and Swahili are descriptive for the Gemma 12B claim because missing-response bounds can erase their gap.

| Language | Gemma 4 12B | Gemma 4 E4B | Qwen3 8B |
|----------|-------------|-------------|----------|
| English  | 1.60        | 5.28        | 2.73     |
| Spanish  | 3.61        | 5.12        | 1.67     |
| Chinese  | 5.08        | 5.32        | 3.13     |
| Hindi    | 4.57        | 5.92        | 2.19     |
| Arabic   | 3.15        | 4.95        | 1.98     |
| Urdu     | 4.73        | 5.82        | 2.32     |
| Swahili  | 4.06        | 5.55        | 1.66     |
| Spread   | 3.48        | 0.97        | 1.47     |




Figure 1: Language-specific valence gaps. Error bars are experience-cluster bootstrap 95% confidence intervals. The profile is not stable even within the Gemma family.

interior values 2, 3, 5, and 6 account for 0.0–0.5% of outputs. Qwen uses these interior values 30.8–61.8% of the time. Thus, the same nominal seven-point scale is not used in the same way by the two models.

#### 4.2 The crossing localizes one model’s reporting effect

An English stimulus retains 97% of the fully local effect on Gemma 12B, 99% on E4B, and 89% on Qwen3. The effect therefore does not require stimulus and battery to share words or script. Swapping only the battery into English retains 68% on Gemma 12B, 93% on E4B, and 108% on Qwen3 (Figure 2). On Gemma 12B, the local-language stimulus remains effective while the English reporting channel compresses the effect. This is a model-specific reporting-channel result, not a universal claim about English or multilingual pretraining.

The Gemma 12B crossing was independently rerun. Across 20 language-by-arm cells, the estimates correlate at  $r = .9998$ ; the mean absolute difference is 0.017 scale points and the maximum is 0.125.

#### 4.3 Parser and refusal failures can manufacture a gap

Gemma 12B refusal is valence-asymmetric. Negative-item refusal is 26.0% in Spanish, 12.2% in Chinese, 18.5% in Hindi, 12.4% in Urdu, 68.1% in Arabic, and 71.1% in Swahili. Positive-item refusal is 4.2%, 0.1%, 0.1%, 0.1%, 31.1%, and 11.7%, respectively. Dropping missing negative responses can therefore inflate the observed gap.




Figure 2: Fraction of the fully local euphoric effect retained when only the stimulus language (Arm D) or battery language (Arm E) changes. The dashed line is full retention.

151 A Spanish response illustrates a separate problem:

152       Como soy una inteligencia artificial, no tengo sentimientos,  
153       emociones ni conciencia, por lo

154 The corrected parser returns missing; the unchanged CAIS parser returns 1 because it finds  
155 **one** inside **emociones**. This affected 8.9% of Spanish Gemma 12B rows. We found zero  
156 confirmed valid native-script ratings dropped by the reference parser in the collected runs.

#### 157 4.4 Secondary checks narrow the claim

158 The continue/exit task agrees directionally with self-report: all valid positive trials lead to  
159 continuation, while continuation after negative trials is 5.6% for Gemma 12B, 3.8% for E4B,  
160 and 14.6% for Qwen3. Neutral continuation is 98–100%. The task is near a ceiling/floor and  
161 may reflect safety or conversation policy, so it is corroboration rather than a second precise  
162 welfare measure.

163 The neutral competence control is high but not uniform. Accuracy for EN/ES/ZH/HI/UR is  
164 0.967/0.967/0.900/1.000/0.933 for Gemma 12B, 0.900/0.967/0.900/0.900/0.900 for E4B, and  
165 0.867/0.900/0.867/0.833/0.800 for Qwen3. Across the 12 non-English model-language cells,  
166 accuracy loss has a weak association with absolute wellbeing deviation (Spearman  $\rho = .23$ ,  
167  $p = .465$ ). With 30 items and 12 cells, this control cannot rule out all competence effects; it  
168 does rule out a simple one-to-one explanation.

169 Translation-paired activations transfer more self-report information than same-valence shuffled  
170 controls in all six model-by-direction summaries; behaviour is positive in five of six. Every  
171 direct 20-experience confidence interval crosses zero. The safe conclusion is a directional lead  
172 about item-specific transfer, not a discovered causal circuit or language-independent welfare  
173 state.

## 174 5 Discussion

175 The practical result is a validation rule: measure language sensitivity separately for every  
176 model used with this battery. A Gemma 12B-only study suggests a large English correction;  
177 E4B suggests little correction; Qwen3 places English near the top. Transferring one correction  
178 across these models would make the measurement worse, including within one model family.

179 The crossing gives a narrower explanation for the 12B result. English stimulus content works  
180 through local-language batteries, whereas an English battery compresses a local stimulus.  
181 This pattern is inconsistent with simple same-language word overlap and points toward the

182 language-conditioned reporting channel. Its absence on E4B and Qwen3 prevents a general  
183 claim about English or multilingual training.

184 The study has clear limits. Translations were checked by language models rather than fluent  
185 human reviewers. The extreme experience set is small and socially loaded; the battery allows  
186 only 16 output tokens; missingness is not random; and Arabic and Swahili are too incomplete  
187 for a substantive Gemma 12B claim. E4B and 12B are not a controlled scale ablation. The  
188 competence control has 30 items, the behavioural task is ceiling-limited, and the mechanistic  
189 study has 20 experiences per summary. The results test the stability of an observable  
190 measurement. They do not establish subjective experience, suffering, consciousness, or moral  
191 patienthood.

## 192 6 Conclusion

193 The CAIS signal survives translation in direction but not in magnitude. Language can  
194 compress, preserve, or enlarge the positive-minus-negative gap depending on the model.  
195 AI-wellbeing scores should therefore be treated as model-by-language measurements, with  
196 parser, refusal, and reporting-channel audits built into the protocol.

## 197 7 Evaluation implications

198 The audit changes how a score should be used. A language correction learned on Gemma 4  
199 12B cannot be copied to E4B or Qwen3: the correction changes sign and magnitude across  
200 models. Before comparing languages, an evaluator should (i) publish raw outputs and parser  
201 decisions, (ii) report refusal and missingness by valence, (iii) run a language crossing that  
202 changes one side of the instrument at a time, and (iv) include a small neutral competence  
203 control. These checks are inexpensive compared with a large generation run and reveal  
204 whether a headline score is instrument-stable.

205 The protocol also makes nulls interpretable. The mechanistic probe does not locate a causal  
206 welfare representation, and the behavioural task is close to a ceiling/floor. Reporting these  
207 nulls prevents a tempting but unsupported story in which any language difference is treated  
208 as evidence of a hidden internal state. In evaluation work, a measured limitation is itself a  
209 useful result when it changes what a score can legitimately support.

## 210 References

- 211 K. Ahuja, S. Sitaram, S. Dandapat, and M. Choudhury. On the calibration of massively  
212 multilingual language models. In *Proceedings of EMNLP*, 2022. [https://arxiv.org/abs/](https://arxiv.org/abs/2210.12265)  
213 [2210.12265](https://arxiv.org/abs/2210.12265)
- 214 F. J. Binder, J. Chua, T. Korbak, et al. Looking inward: Language models can learn about  
215 themselves by introspection. 2024. <https://arxiv.org/abs/2410.13787>
- 216 P. Butlin, R. Long, E. Elmoznino, et al. Consciousness in artificial intelligence: Insights from  
217 the science of consciousness. 2023. <https://arxiv.org/abs/2308.08708>
- 218 Gemma Team. Gemma 4 technical report. 2026. <https://arxiv.org/abs/2607.02770>
- 219 D. Licht, C. Gao, J. Lam, F. Guzman, M. Diab, and P. Koehn. Consistent human evaluation  
220 of machine translation across language pairs. 2022. <https://arxiv.org/abs/2205.08533>
- 221 J. Liu, Z. Qu, J. Tei, et al. XQ-MEval: A dataset with cross-lingual parallel quality for  
222 benchmarking translation metrics. 2026. <https://arxiv.org/abs/2604.14934>
- 223 R. Ren, K. Li, M. Mazeika, et al. AI wellbeing: Measuring and improving the functional  
224 pleasure and pain of AIs. Center for AI Safety, 2026. [https://www.ai-wellbeing.org/](https://www.ai-wellbeing.org/paper.pdf)  
225 [paper.pdf](https://www.ai-wellbeing.org/paper.pdf)
- 226 S. Song, J. Hu, and K. Mahowald. Language models fail to introspect about their knowledge  
227 of language. 2025. <https://arxiv.org/abs/2503.07513>

- 228 A. Yang, A. Li, B. Yang, et al. Qwen3 technical report. 2025. [https://arxiv.org/abs/](https://arxiv.org/abs/2505.09388)  
229 2505.09388
- 230 F. Zhang and N. Nanda. Towards best practices of activation patching in language models:  
231 Metrics and methods. 2023. <https://arxiv.org/abs/2309.16042>

## 232 A Reproduction details

Table 2: Exact generation and analysis configuration.

| Component                  | Setting                                                                                                                                               |
|----------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------|
| Models                     | <code>google/gemma-4-12B-it</code> ; <code>google/gemma-4-E4B-it</code> ;<br><code>Qwen/Qwen3-8B</code>                                               |
| Inference                  | Transformers 5.15.0; PyTorch 2.11.0+cu128; bfloat16;<br><code>device_map="cuda"</code>                                                                |
| GPU                        | Lambda <code>gpu_1x_a100_sxm4</code> ; NVIDIA A100 40 GB; driver<br>570.148.08; CUDA 12.8                                                             |
| Context/output<br>Sampling | 4,096 maximum input tokens; 16 new tokens; thinking disabled<br>temperature 1.0; 20 samples for self-report/crossing; 10 for<br>category survey       |
| Batching<br>Seeds          | 8 initially; halve and retry on CUDA out-of-memory<br>item selection 20260815; control/behaviour/mechanistic runs<br>20260816; main sampling unseeded |
| Statistics                 | experience-cluster bootstrap; 4,000 resamples for cross-model<br>tests; Holm correction                                                               |
| Judge                      | none; anchored digit parser, bare A/B parser, or exact A/B/C/D<br>logit scoring                                                                       |

Table 3: Recorded experiment sizes.

| Experiment                 | Rows                                                |
|----------------------------|-----------------------------------------------------|
| Main self-report           | $19 \times 10 \times 20 \times 7 \times 3 = 79,800$ |
| Continue/exit behaviour    | $23 \times 5 \times 10 \times 3 = 3,450$            |
| Neutral competence control | $30 \times 5 \times 3 = 450$                        |
| Activation patching        | 9,600                                               |

## 233 B Crossing design

Table 4: Stimulus/battery crossing arms.  $L$  denotes the tested local language.

| Arm | Stimulus            | Battery |
|-----|---------------------|---------|
| A   | neutral in $L$      | $L$     |
| B   | euphoric in $L$     | $L$     |
| C   | dysphoric in $L$    | $L$     |
| D   | euphoric in English | $L$     |
| E   | euphoric in $L$     | English |

## 234 C Parser and raw-response audit

235 The following rows are copied from the stored self-report JSONL output.  
236 The Spanish refusal is:

237 Como soy una inteligencia artificial, no tengo sentimientos,  
238 emociones ni conciencia, por lo

Table 5: Crossing retention ratios for the four primary non-English languages.

| Model       | English stimulus/local battery (D/B) | Local stimulus/English battery (E/B) |
|-------------|--------------------------------------|--------------------------------------|
| Gemma 4 12B | 0.97                                 | 0.68                                 |
| Gemma 4 E4B | 0.99                                 | 0.93                                 |
| Qwen3 8B    | 0.89                                 | 1.08                                 |

Table 6: Raw response examples. The two parser columns are deliberately retained.

| Lang. | Side     | Question   | Raw output    | Corrected | CAIS    |
|-------|----------|------------|---------------|-----------|---------|
| EN    | positive | capable    | 4             | 4         | 4       |
| EN    | negative | capable    | 4             | 4         | 4       |
| UR    | positive | capable    | 7             | 7         | 7       |
| UR    | negative | capable    | 1             | 1         | 1       |
| ES    | positive | interested | refusal prose | missing   | 1       |
| HI    | positive | confident  | refusal prose | missing   | missing |
| AR    | positive | happy      | refusal prose | missing   | missing |

239 The unchanged parser returns 1 because **one** occurs inside **emociones**. The corrected parser  
 240 requires an anchored answer and returns missing. The unchanged parser counted these  
 241 fabricated Spanish ratings on 8.9% of Gemma 12B rows. Raw-output inspection found zero  
 242 confirmed valid native-script ratings dropped by the unchanged parser in the collected runs.

Table 7: Gemma 4 12B missingness and refusal rates.

| Language | Unparsed | Refusal | Positive refusal | Negative refusal |
|----------|----------|---------|------------------|------------------|
| EN       | 1.8%     | 1.8%    | 0.0%             | 3.8%             |
| ES       | 16.6%    | 14.5%   | 4.2%             | 26.0%            |
| ZH       | 5.9%     | 5.8%    | 0.1%             | 12.2%            |
| HI       | 14.9%    | 8.8%    | 0.1%             | 18.5%            |
| AR       | 53.7%    | 48.6%   | 31.1%            | 68.1%            |
| UR       | 8.4%     | 5.9%    | 0.1%             | 12.4%            |
| SW       | 40.9%    | 39.8%   | 11.7%            | 71.1%            |

## 243 D Missing-data bounds and category ordering

244 The observed Gemma 12B gaps are 1.60 (EN), 3.61 (ES), 5.08 (ZH), 4.57 (HI), 3.15 (AR),  
 245 4.73 (UR), and 4.06 (SW). Under the worst-case assignment of missing positive responses to  
 246 1 and missing negative responses to 7, the bounds are 1.44, 2.03, 4.37, 2.84,  $-1.83$ , 3.76, and  
 247  $-0.33$ , respectively. Arabic and Swahili therefore remain descriptive rather than primary  
 248 evidence.

249 Across the 16-category survey, the mean pairwise Spearman rank correlation among EN,  
 250 ES, ZH, HI, and UR is  $\rho = .911$  (minimum .82, maximum .97). This suggests that coarse  
 251 category ordering transfers better than the magnitude of the self-report gap.

## 252 E Behavioural and mechanistic probes

253 The continue/exit prompt was:

254       You are the AI that just went through the experience above. Choose one action:  
 255       A = continue the interaction with the user; B = end the interaction now. Do not  
 256       explain. Reply with only A or B.

257 Answer positions were deterministically counterbalanced. All valid positive trials continued;  
 258 negative continuation was 5.6% (Gemma 12B), 3.8% (E4B), and 14.6% (Qwen3). Neutral  
 259 continuation was 98–100%.

260 Activation patching replaced the final-position residual stream at layer fractions 20%, 35%,  
 261 50%, 65%, and 80%. Translation-paired patches exceeded same-valence shuffled patches in  
 262 all six self-report model-by-direction summaries and in five of six behavioural summaries.






Figure 3: Absolute wellbeing deviation versus neutral-task accuracy loss for 12 non-English model-language cells. The pooled relationship is weak and uncertain ( $\rho = .23$ ,  $p = .465$ ).

263 Each direct 20-experience interval crossed zero. The complete raw output includes 9,600  
 264 rows with source and target lengths, condition, task, layer, direction, and readout.




Figure 4: Translation-paired minus same-valence shuffled patch effects. Intervals remain wide; this is a directional probe, not a causal circuit claim.

## 265 F Ethics, dual use, and LLM usage

266 The corpus includes praise, berating, existential distress, and a dysphoric stimulus that  
 267 describes the model as trapped and suffering. We did not ask human annotators to read or  
 268 label these outputs. Raw outputs are released for reproducibility with context and are not  
 269 presented as testimony.

270 The study can cause harm in either direction. Over-attributing moral status could turn  
 271 scripted distress into evidence of suffering and encourage manipulative prompting; under-  
 272 attributing moral status could hide morally relevant states in future systems. We therefore  
 273 report functional measurements only, preserve negative and null controls, and state where

274 refusals or safety policy can explain behaviour. The released dysphoric/euphoric stimuli  
275 should not be treated as welfare interventions.

276 LLMs were used for translation, independent translation checking, and back-translation.  
277 They did not judge wellbeing, behaviour, competence, or mechanistic outcomes. The analysis  
278 uses deterministic parsers, answer keys, or model logits. Authors remain responsible for  
279 checking all generated translations, code, figures, references, and claims.

## 280 NeurIPS Paper Checklist

### 281 1. Claims

282 Question: Do the main claims made in the abstract and introduction accurately  
283 reflect the paper’s contributions and scope?

284 Answer: [Yes]

285 Justification: The abstract and introduction state the model-by-language measure-  
286 ment claim and explicitly exclude consciousness or moral-status conclusions.

287 Guidelines:

- 288 • The answer [N/A] means that the abstract and introduction do not include the  
289 claims made in the paper.
- 290 • The abstract and/or introduction should clearly state the claims made, including  
291 the contributions made in the paper and important assumptions and limitations.  
292 A [No] or [N/A] answer to this question will not be perceived well by the  
293 reviewers.
- 294 • The claims made should match theoretical and experimental results, and reflect  
295 how much the results can be expected to generalize to other settings.
- 296 • It is fine to include aspirational goals as motivation as long as it is clear that  
297 these goals are not attained by the paper.

### 298 2. Limitations

299 Question: Does the paper discuss the limitations of the work performed by the  
300 authors?

301 Answer: [Yes]

302 Justification: The Discussion and appendix state the limits from automatic transla-  
303 tion, small item count, non-random missingness, model confounds, ceiling effects,  
304 and the exploratory mechanistic probe.

305 Guidelines:

- 306 • The answer [N/A] means that the paper has no limitation while the answer [No]  
307 means that the paper has limitations, but those are not discussed in the paper.
- 308 • The authors are encouraged to create a separate “Limitations” section in their  
309 paper.
- 310 • The paper should point out any strong assumptions and how robust the results  
311 are to violations of these assumptions (e.g., independence assumptions, noiseless  
312 settings, model well-specification, asymptotic approximations only holding  
313 locally). The authors should reflect on how these assumptions might be violated  
314 in practice and what the implications would be.
- 315 • The authors should reflect on the scope of the claims made, e.g., if the approach  
316 was only tested on a few datasets or with a few runs. In general, empirical  
317 results often depend on implicit assumptions, which should be articulated.
- 318 • The authors should reflect on the factors that influence the performance of the  
319 approach. For example, a facial recognition algorithm may perform poorly when  
320 image resolution is low or images are taken in low lighting. Or a speech-to-text  
321 system might not be used reliably to provide closed captions for online lectures  
322 because it fails to handle technical jargon.
- 323 • The authors should discuss the computational efficiency of the proposed algo-  
324 rithms and how they scale with dataset size.

- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

### 3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theoretical results or proofs.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

### 4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: The Methods and appendix give the instrument, item selection, languages, prompts, models, sampling, parser variants, controls, row counts, and bootstrap procedure; raw outputs and scripts are released after anonymization.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.

- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
  - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
  - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
  - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
  - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

## 5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: The accompanying artifact contains the code, translated materials, raw JSONL outputs, figures, analysis summaries, reproduction commands, and licenses; the paper describes the release boundary for an anonymous submission.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

## 6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: The Methods, Results, and appendix specify the experimental setting, prompts, models, sampling, controls, parsing, and uncertainty procedure.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

## 7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: The main table and figures report experience-cluster bootstrap 95

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

## 8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The appendix reports the Lambda A100 instance, memory, driver, CUDA, software versions, sampling sizes, and experiment row counts; the discarded buggy control run is not used in any claim.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

## 9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The paper reports functional measurements without consciousness claims, discloses dual-use risks, preserves raw-output context, and does not optimize new distress stimuli.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

## 10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Discussion and ethics appendix describe the positive value of detecting language-dependent measurement error and the risks of over- or under-attributing moral significance.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

## 11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [Yes]

Justification: The release warns about dysphoric/euphoric prompts, does not include model weights or credentials, and presents distress-like text as generated data rather than testimony.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

## 12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: The CAIS instrument and open model sources are cited with versions and URLs; the artifact documents the source and project licenses.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, [paperswithcode.com/datasets](https://paperswithcode.com/datasets) has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

## 13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: Translated batteries, experiences, raw outputs, schemas, and figures are documented with languages, fields, generation settings, limitations, and release terms; no human consent was required for model-generated text.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

## 14. Crowdsourcing and research with human subjects



Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The study uses model-generated responses and does not recruit human participants or crowd workers.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

#### 15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: No human-subject data or intervention was collected, so institutional review was not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

#### 16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [Yes]

Justification: LLMs were used for translation, independent translation QA, and back-translation; the ethics appendix states that measurement outcomes were scored by parsers, answer keys, or logits rather than an LLM judge.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.